

Real-Time Object Detection with OpenCV

Represented by : Basma Ahmed .



/PixelsEgypt



www.pixelseg.com



/PixelsEgyptOrg



/Company/PixelsEgyptOrg



/Pixels-HU

The content for today:

1. What Are Pre-trained Models and Neural Networks?
2. Face and Object Detection .
3. Object Tracking and Color-Based Tracking .
4. Masking and Drawing on Video.
5. Using Pretrained AI Models in OpenCV.
6. Real-Time Detection with MobileNet-SSD.
7. Hands-On: Live Object Detection with Webcam.


1. Pre-trained Models and Neural Networks

Pre-trained Model : is a model that was already trained by experts on a huge dataset (like millions of images).

- Instead of building a new model from scratch, we **reuse** it to do tasks like recognizing or detecting objects in images or videos.

What is a Neural Network?

- A **neural network** is a system that learns how to recognize patterns for example, what a cat or a car looks like.
- A special type called a **Convolutional Neural Network (CNN)** is very good at understanding **images**.



1.2 Why Do We Use Pre-trained Models?

- **Saves time and effort** — no need to collect and train on huge datasets.
- **Gives high accuracy** — these models already know how to detect many objects.
- **Works well with OpenCV** — you can easily load the model and start detecting objects live from your camera.


2. Face and Object Detection:

1. What is Face Detection?

- The process of identifying and locating **human faces** in digital images or video streams.
- **Not** the same as face recognition (which identifies *who* the person is).

2. Why Use OpenCV?

- OpenCV is an open-source computer vision library with built-in tools for **face detection**.
- Supports both traditional (**Haar**) and deep learning-based methods.



2.2 methods used for face detection :


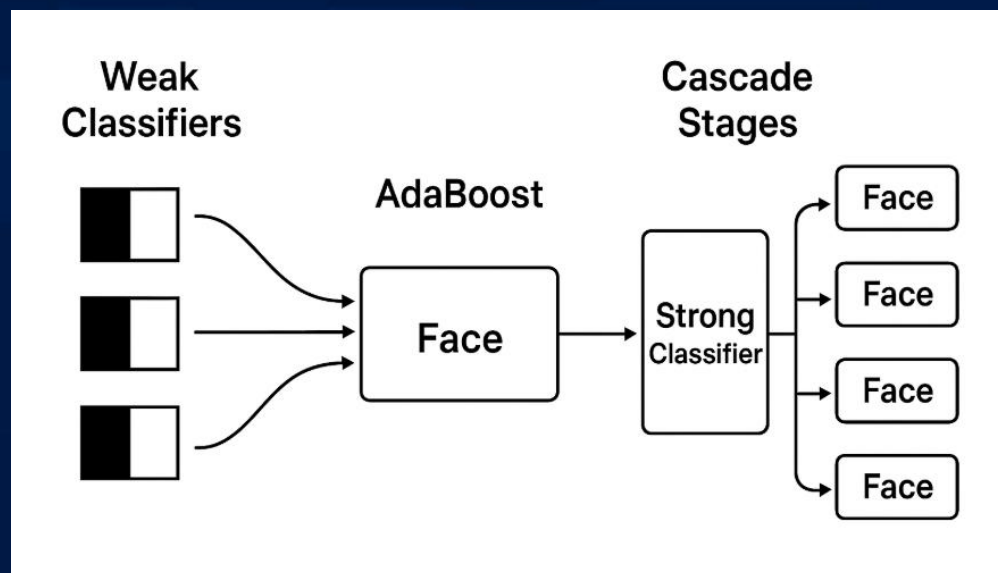

1. Haar Cascade Classifier:

- A **machine learning-based** approach proposed by Viola and Jones.
- Uses **features** like edges, lines, and rectangles.
- Lightweight and **real-time on CPU**.

How it works:

- Uses a series of **Haar-like features** (simple rectangular patterns).
- Applies a **cascade of classifiers** trained via **AdaBoost**.
- Fast because it rejects non-face regions quickly.

Example :



2.2 methods used for face detection :

2. DNN-Based Detection (Deep Neural Networks):

- Caffe-based OpenCV DNN face detector (res10_300x300_ssd_iter_140000.caffemodel).
- MTCNN (Multi-task Cascaded Convolutional Networks).
- RetinaFace, BlazeFace, YOLO-based face detectors.

How it works:

- Uses **deep convolutional neural networks (CNNs)** to learn complex patterns.
- Detects faces more robustly under varied conditions (lighting, angle, size).

2.3 Key Parameters for Face Detection:

1. **scaleFactor:**


Purpose: Controls how the image is resized during detection to handle faces of different sizes.

Example Values:

- **1.1:** Resizes the image by 10% at each scale. Higher accuracy, but slower.
- **1.5:** Resizes by 50% at each scale. Faster but less precise.

Effect: Lower values **detect smaller faces** but take more time.

2.3 Key Parameters for Face Detection:



2. minNeighbors:

Purpose: Specifies how many overlapping rectangles are required to confirm a face.

Example Values:

- 3: Less strict, detects more faces but may include false positives.
- 5: Stricter, fewer false positives but may miss some faces.

Effect: Higher values **reduce false positives** but might miss some faces.

2.3 Key Parameters for Face Detection:

3. minSize :


Purpose: Sets the minimum size (width, height) of the detected face. Smaller objects will be ignored.

Example Values:

- (30, 30): Ignores faces smaller than 30x30 pixels. Helps improve speed and reduce false detections.


Effect: **Filtering out small regions** can make detection faster and more accurate, especially in high-resolution images.

2.4 Setting Up Haar Cascade Classifier :




```
1 import cv2
2
3 # Load the pre-trained Haar cascade classifier for face detection
4 face_cascade = cv2.CascadeClassifier(cv2.data.haarcascades + 'haarcascade_frontalface_default.xml')
5 image = cv2.imread('new.png')
6
7 gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
8
9 # Detect faces in the image
10 faces = face_cascade.detectMultiScale(gray, scaleFactor=1.1, minNeighbors=5, minSize=(30, 30))
11
12 # Draw rectangles around detected faces
13 for (x, y, w, h) in faces:
14     cv2.rectangle(image, (x, y), (x + w, y + h), (255, 0, 0), 3) # Blue rectangle
15
16 cv2.imshow( winname: 'Face Detection', image)
17 cv2.waitKey(0)
18 cv2.destroyAllWindows()
```

The Output :



Task 1:



1. Detect a face from one image **Using OpenCV's Haar Cascade classifier** .
2. Crop it .
3. Resize and rotate it .
4. Paste it onto a white image (like placing a sticker) .




3. Object Tracking and Color-Based Tracking :

1. What is Object Tracking?

- Tracks an object **after it's been detected** in the first frame.
- Doesn't re-detect the object in future frames—just estimates movement based on the previous position.

How to Create a Tracker ?

1. `cv2.legacy.TrackerMIL_create()`
2. `cv2.legacy.TrackerKCF_create()`
3. `cv2.legacy.TrackerTLD_create()`
4. `cv2.legacy.TrackerMOSSE_create()`
5. `cv2.legacy.TrackerCSRT_create()`



3.1 Which OpenCV Tracker Should You Use?

Feature	CSRT	KCF	MOSSE
Accuracy	Very high	Medium to good	Low
Speed	Slower	Faster than CSRT	Fastest
Scale Adaptation	Supports scale changes	Does not support scale	Does not support scale
Handles Deformation	Very good	Average	Poor
Best Use Case	Precise tracking in complex scenes	Balance between speed and accuracy	Real-time, simple tracking
Ideal for	When object size or shape changes	When object is stable	When you need instant speed

3.2 Setting Up Object Tracking:

1.install Required Package : **pip install opencv-contrib-python**

Example :

```
110 # Initialize tracker
111 tracker = cv2.TrackerCSRT_create()
112 tracker.init(frame, bbox)
113
114 # Start tracking
115 while True:
116     ret, frame = cap.read()
117     if not ret:
118         break
119
120     success, bbox = tracker.update(frame)
121
122     if success:
123         x, y, w, h = [int(v) for v in bbox]
124         cv2.rectangle(frame, (x, y), (x+w, y+h), (0, 255, 0), 2)
125         cv2.putText(frame, "Tracking", (20, 50),
126                     cv2.FONT_HERSHEY_SIMPLEX, 0.7, (0, 255, 0), 2)
127     else:
128         cv2.putText(frame, "Lost", (20, 50),
129                     cv2.FONT_HERSHEY_SIMPLEX, 0.7, (0, 0, 255), 2)
130
131     cv2.imshow("Object Tracking", frame)
132
133     if cv2.waitKey(1) & 0xFF == ord("q"):
134         break
135
136 cap.release()
137 cv2.destroyAllWindows()
```


3.3 COLOR-BASED TRACKING (Using HSV Color Filtering):

1. What is Color-Based Tracking?

- Tracks objects **based on a specific color** rather than shape or texture.
- Works well for **bright, distinct-colored objects** like a red ball, green gloves, etc.

2. Why Use HSV Instead of RGB?

- RGB is sensitive to lighting.
- **HSV (Hue, Saturation, Value)** allows better control of color filtering.

Example :

```
234 import cv2
235 import numpy as np
236
237 cap = cv2.VideoCapture(0)
238 while True:
239     ret, frame = cap.read()
240     if not ret:
241         break
242
243     hsv = cv2.cvtColor(frame, cv2.COLOR_BGR2HSV)
244
245     # Track red color
246     lower_red = np.array([0, 120, 70])
247     upper_red = np.array([10, 255, 255])
248     mask = cv2.inRange(hsv, lower_red, upper_red)
249
250     # Find contours
251     contours, _ = cv2.findContours(mask, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
252
253     for contour in contours:
254         if cv2.contourArea(contour) > 300:
255             x, y, w, h = cv2.boundingRect(contour)
256             cv2.rectangle(frame, (x, y), (x+w, y+h), (0, 255, 0), 2)
257             cv2.putText(frame, text="Red Object", org=(x, y - 10),
258                         cv2.FONT_HERSHEY_SIMPLEX, fontScale=0.6, color=(0, 255, 0), thickness=2)
259
260     cv2.imshow(winname="Mask + Drawing", frame)
```

Task 2:

- 1- Track an object of a specific color
(like blue or red) in real-time using your webcam.
- 2- draw a circle around it on the live video feed .



4. Masking and Drawing on Video:

1. Masking :

It is the process of using a **binary image** (mask) to define which parts of an image to show or hide.

➤ A **mask** is usually black and white (**0s and 255s**).

How do we apply it to the image?

We apply a mask to an image using **bitwise operations**, specifically **cv2.bitwise_and()** in OpenCV.

4.1 Bitwise Operations on Images :


Common Operations :

- I. `cv2.bitwise_and(img1, img2)` : Performs **AND** operation .
- I. `cv2.bitwise_or(img1, img2)` : Performs **OR** operation .
- II. `cv2.bitwise_xor(img1, img2)` : Performs **XOR** operation .
- III. `cv2.bitwise_not(img)` : Performs **NOT** operation .

Example on Common Operations :

```
1 import cv2
2
3 img1 = cv2.imread('new.png')
4 img2 = cv2.imread("C:\\Users\\Basma Ahmed\\Pictures\\Screenshots\\Screenshot 2024-12-07 065433.png")
5
6 # Resize both images to the same size
7 img2 = cv2.resize(img2, dsize=(img1.shape[1], img1.shape[0]))
8
9 bitwise_and = cv2.bitwise_and(img1, img2)
10
11 cv2.imshow( winname: 'Image 1', img1)
12 cv2.imshow( winname: 'Image 2', img2)
13 cv2.imshow( winname: 'Bitwise AND', bitwise_and)
14
15 cv2.waitKey(0)
16 cv2.destroyAllWindows()
```

The Output :




Task 3:

Combine two images using `cv2.bitwise_and`, `cv2.bitwise_or`, `cv2.bitwise_xor`, and `cv2.bitwise_not`, and visualize the differences.




4.2 Masking in Real-time Video :




```
161 cap = cv2.VideoCapture(0)
162
163 while True:
164     ret, frame = cap.read()
165     if not ret:
166         break
167
168     # Convert to HSV color space
169     hsv = cv2.cvtColor(frame, cv2.COLOR_BGR2HSV)
170
171     # Define color range for blue
172     lower_blue = np.array([100, 150, 50])
173     upper_blue = np.array([140, 255, 255])
174
175     # Create mask
176     mask = cv2.inRange(hsv, lower_blue, upper_blue)
177
178     # Apply mask
179     masked_output = cv2.bitwise_and(frame, frame, mask=mask)
180
181     # Show both original and masked output
182     cv2.imshow( winname: "Original", frame)
183     cv2.imshow( winname: "Masked (Blue Only)", masked_output)
184
185     if cv2.waitKey(1) & 0xFF == ord('q'):
186         break
187 cap.release()
188 cv2.destroyAllWindows()
```


4.3 Drawing on Video Frames :

```
189 import cv2
190 cap = cv2.VideoCapture(0)
191 if not cap.isOpened():
192     print("Error: Could not open webcam.")
193     exit()
194 x = 0
195 direction = 1
196 while True:
197     ret, frame = cap.read()
198     if not ret:
199         break
200
201     height, width = frame.shape[:2]
202     center = (x, height // 2)
203     cv2.circle(frame, center, radius: 50, color: (255, 0, 0), thickness: 3)
204     cv2.line(frame, pt1: (0, 0), center, color: (0, 0, 255), thickness: 2)
205
206     x += 5 * direction
207     if x >= width or x <= 0:
208         direction *= -1 # change direction
209
210     cv2.imshow( winname: "Drawing on Video", frame)
211
212     if cv2.waitKey(1) & 0xFF == ord('q'):
213         break
214
215 cap.release()
216 cv2.destroyAllWindows()
```



BreakTime



5. Using Pretrained AI Models in OpenCV:

Recap:

A **pretrained model** is a machine learning model that has already been trained on a **large dataset** (like ImageNet, COCO, etc.) and can be reused for similar tasks without retraining.

2. Examples of pretrained models used with OpenCV:

- Face detectors (based on **DNN** or **Haar Cascades**).
- Object detection (e.g., **MobileNet + SSD**).
- Pose estimation.
- Age & gender prediction.

5.2 How does OpenCV use pretrained models?

OpenCV provides a tool called **cv2.dnn** (DNN = Deep Neural Network), which lets us:

1. **Load** the pretrained model into our code.
2. **Feed images or video frames** into it.
3. **Get predictions** like: “There’s a face here, at this location.”

Example:

Think of the pretrained model like a **scanner**:

1. You give it a picture.
2. It looks at it.
3. It tells you what it sees: “I see a face in this part of the image.”

5.3 Face Detection Using Pretrained DNN:

OpenCV supports a pre-trained model called **Res10 SSD Face Detector**, trained using a framework called **Caffe**.

Steps	Description
1. Load the model files	- deploy.prototxt.txt: Defines the model architecture - res10_300x300_ssd_iter_140000.caffemodel: Contains the pre-trained weights
2. Capture a frame	Read a frame from the webcam using cv2.VideoCapture()
3. Convert to blob	Use cv2.dnn.blobFromImage() to preprocess the frame for the neural network
4. Set input to the model	Feed the blob into the model using net.setInput()
5. Forward pass	Run net.forward() to get detection predictions
6. Draw results	Use OpenCV drawing functions (e.g., cv2.rectangle) to display detected faces

Example :

```
258 # Load the pretrained model (Caffe format)
259 model_file = "res10_300x300_ssd_iter_140000.caffemodel"
260 config_file = "deploy.prototxt.txt"
261
262 net = cv2.dnn.readNetFromCaffe(config_file, model_file)
263 # Start video capture from webcam
264 cap = cv2.VideoCapture(0)
265 while True:
266     ret, frame = cap.read()
267     if not ret:
268         break
269
270     # Prepare the frame as input blob for the model
271     blob = cv2.dnn.blobFromImage(frame, scaleFactor: 1.0, size: (300, 300), mean: (104, 117, 123))
272     net.setInput(blob)
273     detections = net.forward()
274     h, w = frame.shape[:2]
```

```
276 # Loop through detections
277 for i in range(detections.shape[2]):
278     confidence = detections[0, 0, i, 2]
279
280     if confidence > 0.5:
281         # Get coordinates of the bounding box
282         box = detections[0, 0, i, 3:7] * [w, h, w, h]
283         x1, y1, x2, y2 = box.astype(int)
284
285         # Draw the bounding box
286         cv2.rectangle(frame, (x1, y1), (x2, y2), (0, 255, 0), 2)
287         label = f"{confidence * 100:.1f}%"
288         cv2.putText(frame, label, org: (x1, y1 - 10),
289                     cv2.FONT_HERSHEY_SIMPLEX, fontScale: 0.6, color: (0, 255, 0), thickness: 2)
290
291 # Display the result
292 cv2.imshow( winname: "Face Detection (DNN)", frame)
293
294 # Press 'q' to quit
295 if cv2.waitKey(1) & 0xFF == ord("q"):
296     break
297
298 cap.release()
299 cv2.destroyAllWindows()
```


6. Real-Time Detection with MobileNet-SSD :

1. What is MobileNet-SSD?

- **MobileNet** is a lightweight, efficient deep learning model — designed to run fast on devices like laptops or even phones.
- **SSD** stands for **Single Shot Multibox Detector** — a type of object detection method that detects multiple objects in a **single pass**, in **real-time**.



6. Real-Time Detection with MobileNet-SSD :

MobileNet-SSD is a model that can **quickly and accurately detect multiple objects** in an image or video stream — like people, cars, bottles, dogs, etc.

Model Files:

- MobileNetSSD_deploy.prototxt → model structure .
- MobileNetSSD_deploy.caffemodel → pretrained weights.

6.2 Steps to Use MobileNet-SSD :

1. Load the model using OpenCV's `cv2.dnn.readNetFromCaffe()`.
2. Capture video frames from webcam or a video.
3. Convert each frame into a blob (input format for the model).
4. Feed the blob into the network using `net.setInput()`.
5. Run a forward pass to get detections using `net.forward()`.
6. Draw boxes and labels for detected objects.

Example :

```

302 import cv2
303 import numpy as np
304
305 # Load class labels MobileNet SSD was trained on
306 classNames = ["background", "aeroplane", "bicycle", "bird", "boat",
307              "bottle", "bus", "car", "cat", "chair", "cow", "diningtable",
308              "dog", "horse", "motorbike", "person", "pottedplant",
309              "sheep", "sofa", "train", "tvmonitor"]
310
311 # Load model files
312 net = cv2.dnn.readNetFromCaffe("MobileNetSSD_deploy.prototxt",
313                               "MobileNetSSD_deploy.caffemodel")
314
315 cap = cv2.VideoCapture(0)
316
317 while True:
318     ret, frame = cap.read()
319     if not ret:
320         break
321
322     h, w = frame.shape[:2]
323
324     # Convert to blob
325     blob = cv2.dnn.blobFromImage(frame, scalefactor: 0.007843, size: (300, 300), mean: 127.5)
326     net.setInput(blob)
327     detections = net.forward()

```

```


329 # Loop through detections
330 for i in range(detections.shape[2]):
331     confidence = detections[0, 0, i, 2]
332     if confidence > 0.5:
333         class_id = int(detections[0, 0, i, 1])
334         class_name = classNames[class_id]
335         box = detections[0, 0, i, 3:7] * np.array([w, h, w, h])
336         x1, y1, x2, y2 = box.astype("int")
337
338         cv2.rectangle(frame, (x1, y1), (x2, y2), (0, 255, 0), 2)
339         cv2.putText(frame, text=f"{class_name}: {confidence*100:.1f}%",
340                    org: (x1, y1 - 10), cv2.FONT_HERSHEY_SIMPLEX, fontScale: 0.5,
341                    color: (0, 255, 0), thickness: 2)
342
343     cv2.imshow( winname: "MobileNet-SSD Real-Time Detection", frame)
344
345     if cv2.waitKey(1) & 0xFF == ord('q'):
346         break
347
348 cap.release()
349 cv2.destroyAllWindows()

```

Hands-On Practice:


Notebook:

https://colab.research.google.com/drive/1pkAAQF3ikl1Iwh_ylFzIQjcCuswW4hSi?usp=sharing



Any questions





THANK YOU .